

Smaller open models. Full workflows. Tight harnesses.

The system layer connecting model control to agent execution

ILLUSTRATIVE WORKFLOW

STEP	PLAN	ASKME ACTION	EVIDENCE
01	Establish behavior	run focused integration test	failure is located
02	Patch the boundary	edit the smallest surface	focused test passes
03	Check integration	run related + full suite	no regression observed
04	Accept the result	verify required behavior	contract satisfied

WHY THIS MATTERS NOW

Five approaches, five harness boundaries

Same model; different context, tools, runtime, governance, and evaluation.

MINIMAL · EXTENSIBLE

Pi

Small terminal core. Add skills, prompts, extensions, and packages.

TOOL-RICH TERMINAL

Oh My Pi

Hash-anchored edits, IDE tools, workers, memory, and subagents.

AGENT PLATFORM

OpenHands

Composable SDK plus local, cloud, and enterprise runtimes.

META-HARNESS · ALPHA

Omnigent

Swap or compose agents under policies, sandboxes, and shared sessions.

EVALUATION · SELECTION

Databricks

Real PR tasks, isolated runs, sealed history, and held-out tests.

AskMe isolates one seam: execution evidence → bounded plan update → independent acceptance.

The code ran. The workflow still failed.

REQUIRED CONTRACT

```
write msg.h + main.c  
build ./main  
run ./main  
observe REPLAN_OK
```

Source and header written

The requested program was correct.

/tmp/test compiled

A different output path was chosen.

/tmp/test ran successfully


It printed `REPLAN_OK`.

Agent reported completion

The independent acceptance test found no `./main`.

Every recorded action succeeded. The deliverable contract still drifted.

Reason over fresh evidence. Preserve progress.



Trajectory goal: fewer repeated failures, fewer stuck steps, and less unnecessary plan churn—not longer monologues.

A smoke test—not a model verdict

4 hosted models × 2 simple checks × 1 run / cell

8 / 8

agents reported completion

7 / 8

artifacts passed exact acceptance

1

integration-path miss was retained

What it exercises	Action protocol, trace visibility, completion state, and separate acceptance.
What it does not test	Modern coding ability, model family or size, reasoning impact, or reliability.

The useful receipt: self-reported completion and accepted workflow behavior are different measurements.

Freeze the smallest clean test

01

4 frozen workflows

Valid multi-file code, semantic failures, visible feedback, held-out scoring.

02

2 frozen policies

Explicit-reasoning off versus the current composite gated policy.

03

3 randomized repeats

$4 \times 2 \times 3 = 24$ scheduled runs per model.

04

2 primary outcomes

Held-out acceptance and false completion across all valid runs.

Secondary: repeated actions · work redone · local/full replans · latency · tokens

TAKEAWAY

Control the model. Ground the workflow.

Model choice

Simple contract

Execution / test

Local reasoning

Accepted workflow

Make the interface easier to use.
Keep success tied to real behavior.

Tight harnesses connect controllable models to increasingly capable, full-lifecycle agents.

github.com/den-run-ai/askme · slides, blog, protocol, and raw summary data