

Smaller open models. Full workflows. Tight harnesses.

The system layer connecting model control to agent execution

ILLUSTRATIVE WORKFLOW

STEP	PLAN	ASKME ACTION	EVIDENCE
01	Establish behavior	run focused integration test	failure is located
02	Patch the boundary	edit the smallest surface	focused test passes
03	Check integration	run related + full suite	no regression observed
04	Accept the result	verify required behavior	contract satisfied

Two trends meet at the harness

SMALLER, OPEN MODELS

More of the model stack is yours

- Execution speed and cost
- Hardware and deployment choice
- Private, edge, or enterprise runtime
- Direct post-training access
- Potential builder and talent flywheel

THE BRIDGE

Tight harnesses

Context · tools · memory · permissions
· workflow · evaluation

FULL-WORKFLOW AGENTS

More of the lifecycle is agentic

- Plan and coordinate
- Search, edit, and execute
- Test and verify
- Preserve state across handoffs
- Improve workflows—and agents

Model control expands what you can deploy. Harness design shapes how the system executes and verifies completion.

The code ran. The workflow still failed.

REQUIRED CONTRACT

```
write msg.h + main.c  
build ./main  
run ./main  
observe REPLAN_OK
```

Source and header written

The requested program was correct.

`/tmp/test` compiled

A different output path was chosen.

`/tmp/test` ran successfully


It printed `REPLAN_OK`.

Agent reported completion

The independent acceptance test found no `./main`.

Every recorded action succeeded. The deliverable contract still drifted.

Reason over fresh evidence. Preserve progress.



Trajectory goal: fewer repeated failures, fewer stuck steps, and less unnecessary plan churn—not longer monologues.

A smoke test—not a model verdict

4 hosted models × 2 simple checks × 1 run / cell

8 / 8

agents reported completion

7 / 8

artifacts passed exact acceptance

1

integration-path miss was retained

What it exercises	Action protocol, trace visibility, completion state, and separate acceptance.
What it does not test	Modern coding ability, model family or size, reasoning impact, or reliability.

The useful receipt: self-reported completion and accepted workflow behavior are different measurements.

THE NEXT EXPERIMENT

Test the reasoning claim directly

01

Realistic failure

Start with syntactically valid multi-file code and a semantic integration bug.

02

Feedback inside the loop

Return focused execution and test results for a bounded repair.

03

Held-out acceptance

Keep final scoring separate from feedback the agent sees.

04

Matched variants

Repeat and randomize reasoning off, gated, and always-on policies.

Measure: accepted behavior · recovery effort · plan stability

TAKEAWAY

Control the model. Ground the workflow.

Model choice

Simple contract

Execution / test

Local reasoning

Accepted workflow

Make the interface easier to use.
Keep success tied to real behavior.

Tight harnesses connect controllable models to increasingly capable, full-lifecycle agents.

github.com/den-run-ai/askme · slides, blog, protocol, and raw summary data